

ECGR 4105 - HW2 - Source Code

September 27, 2024

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

from sklearn.preprocessing import MinMaxScaler
```

```
[4]: file_url = 'https://raw.githubusercontent.com/lnguye782/ECGR-4105-Intro-to-ML/
↪refs/heads/main/HW2/Housing.csv'
data = pd.read_csv(file_url)

data.head()
```

```
[4]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

```
[5]: # Function for gradient descent for linear regression
def gradient_descent(x, y, learning_rate, iterations):
    m = x.shape[0]
    n = x.shape[1]
    theta = np.zeros(n) # Initialize theta to zero
    b = 0 # Intercept
    losses = []
```

```

for i in range(iterations):
    # Prediction for y
    y_prediction = np.dot(x, theta) + b

    # Calculate Mean Squared Error (loss)
    loss = np.mean((y_prediction - y) ** 2)
    losses.append(loss)

    # Gradient descent update rules
    d_theta = (1/m) * np.dot(x.T, (y_prediction - y))
    d_b = (1/m) * np.sum(y_prediction - y)

    # Update theta and b
    theta -= learning_rate * d_theta
    b -= learning_rate * d_b

return theta, b, losses

```

```

[6]: # Problem 1a
x_a = data[['area', 'bedrooms', 'bathrooms', 'stories', 'parking']]
y_a = data['price']

# Splitting data into 80% and 20% split between training and evaluation
x_train_a, x_eval_a, y_train_a, y_eval_a = train_test_split(x_a, y_a,
    ↪test_size=0.2, random_state=42)

scaler = StandardScaler()
x_train_a_scaled = scaler.fit_transform(x_train_a)
x_eval_a_scaled = scaler.transform(x_eval_a)

```

```

[12]: # Set parameters for gradient descent
learning_rate = 0.01
iterations = 1000

# Training the gradient descent models
theta_train_a, b_train_a, loss_train_a = gradient_descent(x_train_a_scaled,
    ↪y_train_a, learning_rate, iterations)
theta_eval_a, b_eval_a, loss_eval_a = gradient_descent(x_eval_a_scaled,
    ↪y_eval_a, learning_rate, iterations)


```

```

[13]: # Plot the training and validation losses
plt.plot(range(iterations), loss_train_a, label='P1a. Training Loss')
plt.plot(range(iterations), loss_eval_a, label='P1a. Validation Loss',
    ↪linestyle='--')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('P1a. Training and Validation Losses Over Iterations')

```

```
plt.legend()
plt.show()
```



```
[14]: # Problem 1b
x_b = data[['area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
            'guestroom', 'basement', 'hotwaterheating',
            'airconditioning', 'parking', 'prefarea']]
y_b = data['price']

# Converting categorical variables to numeric using pd.get_dummies
x_b = pd.get_dummies(x_b, drop_first=True)

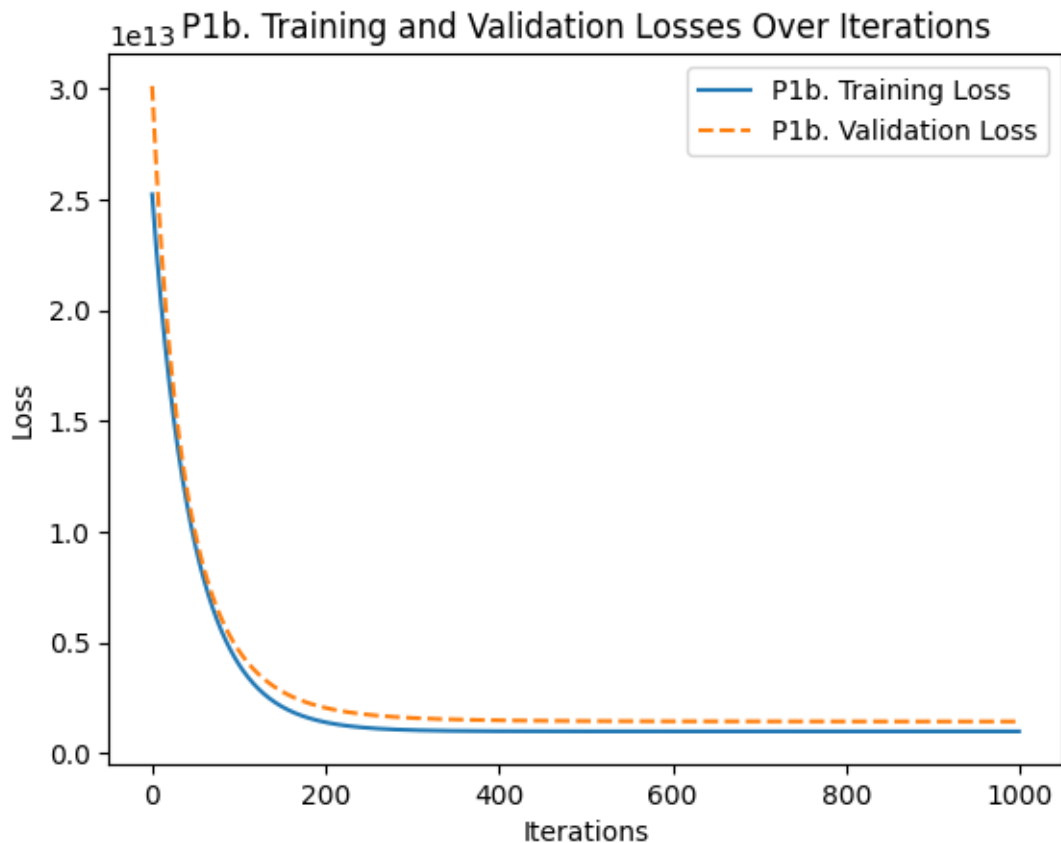
x_train_b, x_eval_b, y_train_b, y_eval_b = train_test_split(x_b, y_b,
    ↪test_size=0.2, random_state=42)

scaler = StandardScaler()
x_train_b_scaled = scaler.fit_transform(x_train_b)
x_eval_b_scaled = scaler.transform(x_eval_b)
```

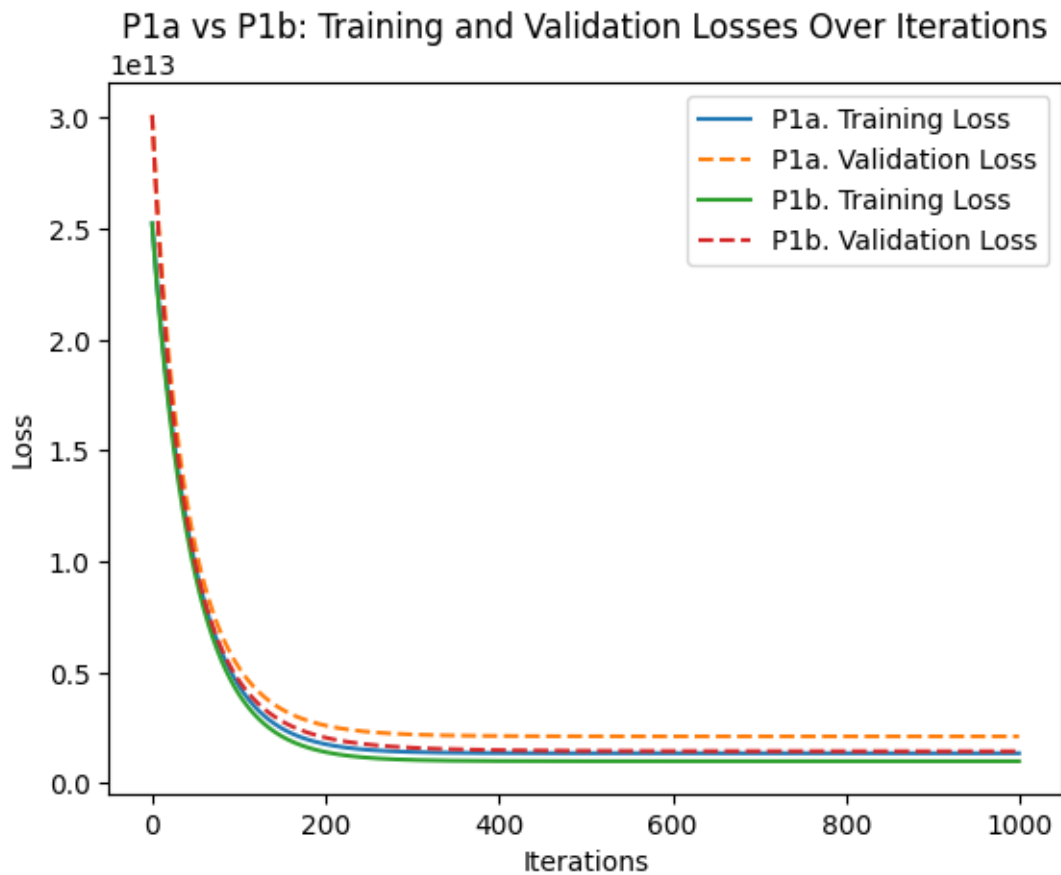
```
[15]: # Set parameters for gradient descent
learning_rate = 0.01
iterations = 1000

# Training the gradient descent models
theta_train_b, b_train_b, loss_train_b = gradient_descent(x_train_b_scaled,
    ↪ y_train_b, learning_rate, iterations)
theta_eval_b, b_eval_b, loss_eval_b = gradient_descent(x_eval_b_scaled,
    ↪ y_eval_b, learning_rate, iterations)
```

```
[16]: plt.plot(range(iterations), loss_train_b, label='P1b. Training Loss')
plt.plot(range(iterations), loss_eval_b, label='P1b. Validation Loss',
    ↪ linestyle='--')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('P1b. Training and Validation Losses Over Iterations')
plt.legend()
plt.show()
```



```
[17]: plt.plot(range(iterations), loss_train_a, label='P1a. Training Loss')
plt.plot(range(iterations), loss_eval_a, label='P1a. Validation Loss',
         linestyle='--')
plt.plot(range(iterations), loss_train_b, label='P1b. Training Loss')
plt.plot(range(iterations), loss_eval_b, label='P1b. Validation Loss',
         linestyle='--')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('P1a vs P1b: Training and Validation Losses Over Iterations')
plt.legend()
plt.show()
```



```
[18]: # Problem 2a
input_variables_a = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking']
x_a = data[input_variables_a]
```

```
[19]: # Function to split and preprocess logic
def preprocess_data_a(scaling_type='standardization', variable=x_a):
    x = variable
```

```

y = data['price']

# Splitting data into 80% and 20% split between training and evaluation
x_train, x_eval, y_train, y_eval = train_test_split(x, y, test_size=0.
↪2, random_state=42)

if scaling_type == 'standardization':
    scaler = StandardScaler()
elif scaling_type == 'normalization':
    scaler = MinMaxScaler()

x_train_scaled = scaler.fit_transform(x_train)
x_eval_scaled = scaler.transform(x_eval)

return x_train_scaled, x_eval_scaled, y_train, y_eval

```

```

[20]: # Function to train model with preprocessing
def train_model_with_preprocessing_a(scaling_type, learning_rate, iterations):
    x_trained_scaled, x_eval_scaled, y_train, y_eval = ↪
    ↪preprocess_data_a(scaling_type=scaling_type)

    theta, b, losses_train = gradient_descent(x_trained_scaled, y_train, ↪
    ↪learning_rate, iterations)

    losses_eval = []
    for i in range(iterations):
        y_pred_eval_iter = np.dot(x_eval_scaled, theta) + b
        loss_eval = (1/(2*x_eval_scaled.shape[0])) * np.sum((y_pred_eval_iter - ↪
        ↪y_eval) ** 2)
        losses_eval.append(loss_eval)

    return theta, b, losses_train, losses_eval

```

```

[21]: # Set parameters
learning_rate = 0.01
iterations = 1000

# Input standardization
theta_std_a, b_std_a, losses_train_std_a, losses_eval_std_a = ↪
    ↪train_model_with_preprocessing_a('standardization', learning_rate, ↪
    ↪iterations)

# Input normalization
theta_norm_a, b_norm_a, losses_train_norm_a, losses_eval_norm_a = ↪
    ↪train_model_with_preprocessing_a('normalization', learning_rate, iterations)

```

```


# Plot the training vs validation losses for input standardization and input
↪normalization
plt.plot(range(iterations), losses_train_std_a, label='Training Loss for
↪Standardization')
plt.plot(range(iterations), losses_eval_std_a, label='Validation Loss for
↪Standardization', linestyle='--')

plt.plot(range(iterations), losses_train_norm_a, label='Training Loss for
↪Normalization')
plt.plot(range(iterations), losses_eval_norm_a, label='Validation Loss for
↪Normalization', linestyle='--')

plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('P2a. Training and Validation Losses for Standardization and
↪Normalization')
plt.legend()
plt.show()

```

P2a. Training and Validation Losses for Standardization and Normalization



```
[22]: # Problem 2b
input_variables_b = ['area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
                    'guestroom', 'basement', 'hotwaterheating',
                    ↪ 'airconditioning',
                    'parking', 'prefarea']
# Converting categorical variables to numeric using pd.get_dummies
x_b = pd.get_dummies(data[input_variables_b], drop_first=True)

[23]: # Function to split and preprocess logic
def preprocess_data_b(scaling_type='standardization', variable=x_b):
    x = variable
    y = data['price']

    # Splitting data into 80% and 20% split between training and evaluation
    x_train, x_eval, y_train, y_eval = train_test_split(x, y, test_size=0.
    ↪ 2, random_state=42)

    if scaling_type == 'standardization':
        scaler = StandardScaler()
    elif scaling_type == 'normalization':
        scaler = MinMaxScaler()

    x_train_scaled = scaler.fit_transform(x_train)
    x_eval_scaled = scaler.transform(x_eval)

    return x_train_scaled, x_eval_scaled, y_train, y_eval

[24]: # Function to train model with preprocessing
def train_model_with_preprocessing_b(scaling_type, learning_rate, iterations):
    x_trained_scaled, x_eval_scaled, y_train, y_eval = ↪
    ↪ preprocess_data_b(scaling_type=scaling_type)

    theta, b, losses_train = gradient_descent(x_trained_scaled, y_train, ↪
    ↪ learning_rate, iterations)

    losses_eval = []
    for i in range(iterations):
        y_pred_eval_iter = np.dot(x_eval_scaled, theta) + b
        loss_eval = (1/(2*x_eval_scaled.shape[0])) * np.sum((y_pred_eval_iter - ↪
    ↪ y_eval) ** 2)
        losses_eval.append(loss_eval)

    return theta, b, losses_train, losses_eval

[25]: # Set parameters
learning_rate = 0.01
iterations = 1000
```



```

# Input standardization
theta_std_b, b_std_b, losses_train_std_b, losses_eval_std_b =
    ↪train_model_with_preprocessing_b('standardization', learning_rate,
    ↪iterations)

# Input normalization
theta_norm_b, b_norm_b, losses_train_norm_b, losses_eval_norm_b =
    ↪train_model_with_preprocessing_b('normalization', learning_rate, iterations)


# Plot the training vs validation losses for input standardization and input
    ↪normalization
plt.plot(range(iterations), losses_train_std_b, label='Training Loss for
    ↪Standardization')
plt.plot(range(iterations), losses_eval_std_b, label='Validation Loss for
    ↪Standardization', linestyle='--')

plt.plot(range(iterations), losses_train_norm_b, label='Training Loss for
    ↪Normalization')
plt.plot(range(iterations), losses_eval_norm_b, label='Validation Loss for
    ↪Normalization', linestyle='--')

plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('P2b. Training and Validation Losses for Standardization and
    ↪Normalization')
plt.legend()
plt.show()

```

P2b. Training and Validation Losses for Standardization and Normalization



```
[29]: # Function for gradient descent with L2 regularization
def gradient_descent_with_regularization(x, y, learning_rate, iterations,
    ↪ lambda_reg=0.1):
    m = x.shape[0]
    n = x.shape[1]
    theta = np.zeros(n) # Initialize theta to zero
    b = 0 # Intercept
    losses = []

    for i in range(iterations):
        # Prediction for y
        y_prediction = np.dot(x, theta) + b

        # Calculate Mean Squared Error (loss)
        loss = (1/(2*m)) * (np.sum((y_prediction - y) ** 2) + lambda_reg * np.
    ↪ sum(theta ** 2))
        losses.append(loss)

        # Gradient descent update rules
        d_theta = (1/m) * (np.dot(x.T, (y_prediction - y)) + lambda_reg * theta)
        d_b = (1/m) * np.sum(y_prediction - y)
```

```

    # Update theta and b
    theta -= learning_rate * d_theta
    b -= learning_rate * d_b

    return theta, b, losses

```

```

[30]: # Problem 3a
# Function to train model with regularization
def train_model_with_regularization_a(scaling_type, learning_rate, iterations,
    ↪lambda_reg):
    x_trained_scaled, x_eval_scaled, y_train, y_eval =
    ↪preprocess_data_a(scaling_type=scaling_type)

    theta, b, losses_train =
    ↪gradient_descent_with_regularization(x_trained_scaled, y_train,
    ↪learning_rate, iterations, lambda_reg)

    losses_eval = []
    for i in range(iterations):
        y_pred_eval_iter = np.dot(x_eval_scaled, theta) + b
        loss_eval = (1/(2*x_eval_scaled.shape[0])) * np.sum((y_pred_eval_iter -
    ↪y_eval) ** 2)
        losses_eval.append(loss_eval)

    return theta, b, losses_train, losses_eval

```

```

[28]: # Set parameters
learning_rate = 0.01
iterations = 1000
lambda_reg = 0.1


# Train the model with regularization using standardization
theta_reg_a, intercept_reg_a, losses_train_reg_a, losses_eval_reg_a =
    ↪train_model_with_regularization_a('standardization', learning_rate,
    ↪iterations, lambda_reg)

# Plot the training and validation losses
plt.plot(range(iterations), losses_train_reg_a, label='Training Loss with
    ↪Regularization')
plt.plot(range(iterations), losses_eval_reg_a, label='Validation Loss with
    ↪Regularization', linestyle='--')

plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('P3a. Training and Validation Losses with Regularization')

```

```
plt.legend()
plt.show()
```



```
[31]: # Problem 3b
# Function to train model with regularization
def train_model_with_regularization_b(scaling_type, learning_rate, iterations,
    ↪ lambda_reg):
    x_trained_scaled, x_eval_scaled, y_train, y_eval =
    ↪ preprocess_data_b(scaling_type=scaling_type)

    theta, b, losses_train =
    ↪ gradient_descent_with_regularization(x_trained_scaled, y_train,
    ↪ learning_rate, iterations, lambda_reg)

    losses_eval = []
    for i in range(iterations):
        y_pred_eval_iter = np.dot(x_eval_scaled, theta) + b
        loss_eval = (1/(2*x_eval_scaled.shape[0])) * np.sum((y_pred_eval_iter -
    ↪ y_eval) ** 2)
```

```
losses_eval.append(loss_eval)

return theta, b, losses_train, losses_eval
```

```
[34]: # Set parameters
learning_rate = 0.01
iterations = 1000
lambda_reg = 0.1

# Train the model with regularization using standardization
theta_reg_b, intercept_reg_b, losses_train_reg_b, losses_eval_reg_b =
    ↪train_model_with_regularization_b('standardization', learning_rate,
    ↪iterations, lambda_reg)

# Plot the training and validation losses
plt.plot(range(iterations), losses_train_reg_a, label='Training Loss with
    ↪Regularization')
plt.plot(range(iterations), losses_eval_reg_a, label='Validation Loss with
    ↪Regularization', linestyle='--')

plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('P3b. Training and Validation Losses with Regularization')
plt.legend()
plt.show()
```

P3b. Training and Validation Losses with Regularization

