

# ECGR\_4105\_HW5\_Problem\_2\_Source\_Code

November 16, 2024

```
[ ]: # Link to Google Colab: https://colab.research.google.com/drive/1omnTRKEN9-b0JubJq-dpHY5SwLbxL50?usp=sharing
```

```
[6]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
[2]: file_url = 'https://raw.githubusercontent.com/lngyue782/ECGR-4105-Intro-to-ML/refs/heads/main/HW2/Housing.csv'
data = pd.read_csv(file_url)

data.head(), data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null    int64
1   area                 545 non-null    int64
2   bedrooms             545 non-null    int64
3   bathrooms            545 non-null    int64
4   stories              545 non-null    int64
5   mainroad             545 non-null    object
6   guestroom            545 non-null    object
7   basement             545 non-null    object
8   hotwaterheating      545 non-null    object
9   airconditioning      545 non-null    object
10  parking              545 non-null    int64
11  prefarea             545 non-null    object
12  furnishingstatus     545 non-null    object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

```
[2]: (
    price area bedrooms bathrooms stories mainroad guestroom basement \
0 13300000 7420 4 2 3 yes no no
1 12250000 8960 4 4 4 yes no no
2 12250000 9960 3 2 2 yes no yes
3 12215000 7500 4 2 2 yes no yes
4 11410000 7420 4 1 2 yes yes yes

    hotwaterheating airconditioning parking prefarea furnishingstatus
0 no yes 2 yes furnished
1 no yes 3 no furnished
2 no no 2 yes semi-furnished
3 no yes 3 yes furnished
4 no yes 2 no furnished ,
None)
```

```
[3]: # Select relevant columns
selected_columns = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking', 'price']
data_selected = data[selected_columns]

# Features and target variable
X = data_selected.drop(columns='price')
y = data_selected['price']

# Split the data into training (80%) and validation (20%) sets
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features (normalize)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_val_scaled = scaler.transform(X_val)

# Convert target to numpy arrays
y_train = y_train.values
y_val = y_val.values

# Check processed shapes
X_train_scaled.shape, X_val_scaled.shape, y_train.shape, y_val.shape
```

```
[3]: ((436, 5), (109, 5), (436,), (109,))
```

```
[5]: # Linear Regression Model using Gradient Descent
class LinearRegressionGD:
    def __init__(self, learning_rate=0.01, epochs=5000):
        self.learning_rate = learning_rate
        self.epochs = epochs
```

```

        self.weights = None
        self.bias = None

    def initialize_parameters(self, n_features):
        self.weights = np.zeros(n_features)
        self.bias = 0

    def compute_loss(self, y_true, y_pred):
        # Mean Squared Error
        return np.mean((y_true - y_pred) ** 2)

    def train(self, X, y, X_val, y_val):
        n_samples, n_features = X.shape
        self.initialize_parameters(n_features)

        train_loss = []
        val_loss = []

        # Training loop
        for epoch in range(self.epochs):
            # Predictions
            y_pred = np.dot(X, self.weights) + self.bias

            # Compute gradients
            dw = -(2 / n_samples) * np.dot(X.T, (y - y_pred))
            db = -(2 / n_samples) * np.sum(y - y_pred)

            # Update parameters
            self.weights -= self.learning_rate * dw
            self.bias -= self.learning_rate * db

            # Compute loss
            loss = self.compute_loss(y, y_pred)
            train_loss.append(loss)

            # Validation loss
            y_val_pred = np.dot(X_val, self.weights) + self.bias
            val_loss.append(self.compute_loss(y_val, y_val_pred))

            # Print progress every 500 epochs
            if (epoch + 1) % 500 == 0:
                print(f"Epoch {epoch + 1}/{self.epochs}: Train Loss = {loss:.4f}, Validation Loss = {val_loss[-1]:.4f}")

        return train_loss, val_loss

# Train with different learning rates

```

```

learning_rates = [0.1, 0.01, 0.001, 0.0001]
results = {}

for lr in learning_rates:
    print(f"\nTraining with Learning Rate: {lr}")
    model = LinearRegressionGD(learning_rate=lr, epochs=5000)
    train_loss, val_loss = model.train(X_train_scaled, y_train, X_val_scaled,
    ↪y_val)
    results[lr] = (train_loss, val_loss, model.weights, model.bias)

```

Training with Learning Rate: 0.1

```

Epoch 500/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 1000/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 1500/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 2000/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 2500/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 3000/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 3500/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 4000/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 4500/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662
Epoch 5000/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3662

```

Training with Learning Rate: 0.01

```

Epoch 500/5000: Train Loss = 1350009010876.5654, Validation Loss =
2292813178774.1748
Epoch 1000/5000: Train Loss = 1350008211333.5713, Validation Loss =
2292721692029.6787
Epoch 1500/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721546204.0869
Epoch 2000/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545726.9561
Epoch 2500/5000: Train Loss = 1350008211326.5803, Validation Loss =
2292721545725.3770
Epoch 3000/5000: Train Loss = 1350008211326.5801, Validation Loss =
2292721545725.3726
Epoch 3500/5000: Train Loss = 1350008211326.5801, Validation Loss =

```

2292721545725.3726

Epoch 4000/5000: Train Loss = 1350008211326.5801, Validation Loss = 2292721545725.3726

Epoch 4500/5000: Train Loss = 1350008211326.5801, Validation Loss = 2292721545725.3726

Epoch 5000/5000: Train Loss = 1350008211326.5801, Validation Loss = 2292721545725.3726

Training with Learning Rate: 0.001

Epoch 500/5000: Train Loss = 4428910450828.4600, Validation Loss = 6026504766422.3066

Epoch 1000/5000: Train Loss = 1767420758090.9146, Validation Loss = 2894953483806.9399

Epoch 1500/5000: Train Loss = 1407960624252.7695, Validation Loss = 2413169011139.9165

Epoch 2000/5000: Train Loss = 1358342820260.3408, Validation Loss = 2323576806063.2661

Epoch 2500/5000: Train Loss = 1351290148936.9075, Validation Loss = 2302396127424.1353

Epoch 3000/5000: Train Loss = 1350229453600.0400, Validation Loss = 2296156481221.6343

Epoch 3500/5000: Train Loss = 1350052956926.5608, Validation Loss = 2294032466154.8330

Epoch 4000/5000: Train Loss = 1350018866646.1482, Validation Loss = 2293247896566.0298

Epoch 4500/5000: Train Loss = 1350011086718.7402, Validation Loss = 2292942769433.6924

Epoch 5000/5000: Train Loss = 1350009048453.4460, Validation Loss = 2292818935060.5659

Training with Learning Rate: 0.0001

Epoch 500/5000: Train Loss = 20705542685265.9922, Validation Loss = 24839773808034.9570

Epoch 1000/5000: Train Loss = 17056916524493.8281, Validation Loss = 20604240673950.2227

Epoch 1500/5000: Train Loss = 14115695294628.7383, Validation Loss = 17201089604829.7383

Epoch 2000/5000: Train Loss = 11738794306343.0312, Validation Loss = 14458005933892.8828

Epoch 2500/5000: Train Loss = 9813833385830.5566, Validation Loss = 12240607727733.4473

Epoch 3000/5000: Train Loss = 8252043228973.9297, Validation Loss = 10443501933882.5371

Epoch 3500/5000: Train Loss = 6982947086155.4678, Validation Loss = 8983600956273.5117

Epoch 4000/5000: Train Loss = 5950339148635.6992, Validation Loss = 7795091394560.7373

Epoch 4500/5000: Train Loss = 5109218014923.4170, Validation Loss =

6825618923305.2949

Epoch 5000/5000: Train Loss = 4423429686043.9248, Validation Loss = 6033374860840.7676

```
[8]: # Plot Training and Validation Loss for each learning rate
for lr, (train_loss, val_loss, _, _) in results.items():
    plt.figure(figsize=(8, 5))
    plt.plot(range(1, len(train_loss) + 1), train_loss, label="Training Loss")
    plt.plot(range(1, len(val_loss) + 1), val_loss, label="Validation Loss")
    plt.title(f"Loss vs Epochs for Learning Rate = {lr}")
    plt.xlabel("Epochs")
    plt.ylabel("Loss (MSE)")
    plt.legend()
    plt.grid()
    plt.show()
```





