

## ECGR\_4105\_HW7\_Problem\_2

December 10, 2024

```
[1]: import matplotlib.pyplot as plt
import time
import tensorflow as tf

from tensorflow.keras import models, layers
from tensorflow.keras.datasets import cifar10
```

```
[2]: # Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
# Flatten labels
y_train = y_train.flatten()
y_test = y_test.flatten()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 13s  
0us/step

```
[3]: # Define a residual block
def residual_block(input_tensor, filters, stride=1):
    x = layers.Conv2D(filters, (3, 3), strides=stride, padding="same",
        ↪activation="relu")(input_tensor)
    x = layers.BatchNormalization()(x)
    x = layers.Conv2D(filters, (3, 3), strides=1, padding="same")(x)
    x = layers.BatchNormalization()(x)

    shortcut = input_tensor
    if stride != 1 or input_tensor.shape[-1] != filters:
        shortcut = layers.Conv2D(filters, (1, 1), strides=stride,
        ↪padding="same")(input_tensor)

    x = layers.add([x, shortcut])
    x = layers.Activation("relu")(x)
    return x
```

```
[4]: # Build the ResNet-10 model
def build_resnet10(input_shape=(32, 32, 3), num_classes=10):
    inputs = layers.Input(shape=input_shape)
    x = layers.Conv2D(64, (3, 3), strides=1, padding="same",
    ↪activation="relu")(inputs)
    x = layers.BatchNormalization()(x)

    # Add 10 residual blocks
    for _ in range(5):
        x = residual_block(x, 64)

    for _ in range(5):
        x = residual_block(x, 128, stride=2 if _ == 0 else 1)

    x = layers.GlobalAveragePooling2D()(x)
    outputs = layers.Dense(num_classes, activation="softmax")(x)
    return models.Model(inputs, outputs)

[5]: # Initialize and compile ResNet-10
resnet10 = build_resnet10()
resnet10.compile(optimizer="adam", loss="sparse_categorical_crossentropy",
    ↪metrics=["accuracy"])

# Train the model
start_time = time.time()
# history = resnet10.fit(x_train, y_train, epochs=200, validation_data=(x_test,
    ↪y_test))
history = resnet10.fit(x_train, y_train, epochs=20, validation_data=(x_test,
    ↪y_test))
training_time = time.time() - start_time

# Evaluate the model
final_loss, final_accuracy = resnet10.evaluate(x_test, y_test, verbose=0)

# Print results
print(f"Training Time: {training_time:.2f} seconds")
print(f"Final Training Loss: {history.history['loss'][-1]:.4f}")
print(f"Final Validation Accuracy: {final_accuracy:.2f}")
```

```
Epoch 1/20
1563/1563          95s 47ms/step -
accuracy: 0.4550 - loss: 1.5227 - val_accuracy: 0.5373 - val_loss: 1.3772
Epoch 2/20
1563/1563          122s 41ms/step
- accuracy: 0.7033 - loss: 0.8440 - val_accuracy: 0.6642 - val_loss: 1.0881
Epoch 3/20
1563/1563          82s 41ms/step -
```

accuracy: 0.7814 - loss: 0.6301 - val\_accuracy: 0.6376 - val\_loss: 1.0900  
 Epoch 4/20  
 1563/1563 82s 41ms/step -  
 accuracy: 0.8224 - loss: 0.5155 - val\_accuracy: 0.7477 - val\_loss: 0.7354  
 Epoch 5/20  
 1563/1563 63s 40ms/step -  
 accuracy: 0.8543 - loss: 0.4244 - val\_accuracy: 0.8176 - val\_loss: 0.5491  
 Epoch 6/20  
 1563/1563 82s 40ms/step -  
 accuracy: 0.8821 - loss: 0.3473 - val\_accuracy: 0.7975 - val\_loss: 0.6262  
 Epoch 7/20  
 1563/1563 84s 41ms/step -  
 accuracy: 0.9010 - loss: 0.2852 - val\_accuracy: 0.7521 - val\_loss: 0.8851  
 Epoch 8/20  
 1563/1563 83s 42ms/step -  
 accuracy: 0.9179 - loss: 0.2372 - val\_accuracy: 0.8211 - val\_loss: 0.5525  
 Epoch 9/20  
 1563/1563 80s 41ms/step -  
 accuracy: 0.9307 - loss: 0.2032 - val\_accuracy: 0.8185 - val\_loss: 0.5812  
 Epoch 10/20  
 1563/1563 83s 42ms/step -  
 accuracy: 0.9452 - loss: 0.1605 - val\_accuracy: 0.8232 - val\_loss: 0.6071  
 Epoch 11/20  
 1563/1563 81s 41ms/step -  
 accuracy: 0.9557 - loss: 0.1311 - val\_accuracy: 0.8080 - val\_loss: 0.7393  
 Epoch 12/20  
 1563/1563 81s 41ms/step -  
 accuracy: 0.9589 - loss: 0.1159 - val\_accuracy: 0.8147 - val\_loss: 0.7551  
 Epoch 13/20  
 1563/1563 82s 40ms/step -  
 accuracy: 0.9672 - loss: 0.0970 - val\_accuracy: 0.8313 - val\_loss: 0.6567  
 Epoch 14/20  
 1563/1563 83s 41ms/step -  
 accuracy: 0.9663 - loss: 0.0966 - val\_accuracy: 0.8279 - val\_loss: 0.6924  
 Epoch 15/20  
 1563/1563 81s 41ms/step -  
 accuracy: 0.9734 - loss: 0.0781 - val\_accuracy: 0.8252 - val\_loss: 0.7335  
 Epoch 16/20  
 1563/1563 84s 42ms/step -  
 accuracy: 0.9712 - loss: 0.0828 - val\_accuracy: 0.8401 - val\_loss: 0.6740  
 Epoch 17/20  
 1563/1563 80s 40ms/step -  
 accuracy: 0.9787 - loss: 0.0614 - val\_accuracy: 0.8491 - val\_loss: 0.6500  
 Epoch 18/20  
 1563/1563 82s 41ms/step -  
 accuracy: 0.9772 - loss: 0.0671 - val\_accuracy: 0.8490 - val\_loss: 0.6481  
 Epoch 19/20  
 1563/1563 83s 42ms/step -

accuracy: 0.9809 - loss: 0.0554 - val\_accuracy: 0.8425 - val\_loss: 0.6675  
Epoch 20/20  
1563/1563 82s 42ms/step -  
accuracy: 0.9773 - loss: 0.0638 - val\_accuracy: 0.8504 - val\_loss: 0.6667  
Training Time: 1694.00 seconds  
Final Training Loss: 0.0694  
Final Validation Accuracy: 0.85

```
[6]: # Compare performance
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")
plt.plot(history.history["accuracy"], label="Training Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title("ResNet-10 Training vs Validation Accuracy")
plt.legend()
plt.show()
```

